



Neural Optimization with Adaptive Heuristics for Intelligent Marketing System

Changshuai Wei*
LinkedIn Corporation
Seattle, USA
chawei@linkedin.com

Benjamin Zelditch
LinkedIn Corporation
New York, USA
bzelditch@linkedin.com

Joyce Chen
LinkedIn Corporation
Sunnyvale, USA
joychen@linkedin.com

Andre Assuncao Silva T Ribeiro
LinkedIn Corporation
Sunnyvale, USA
aribeiro@linkedin.com

Jingyi Kenneth Tay
LinkedIn Corporation
Sunnyvale, USA
ktay@linkedin.com

Borja Ocejo Elizondo
LinkedIn Corporation
Sunnyvale, USA
bocejo@linkedin.com

Sathiya Keerthi Selvaraj
LinkedIn Corporation
Sunnyvale, USA
keselvaraj@linkedin.com

Aman Gupta
LinkedIn Corporation
Sunnyvale, USA
amagupta@linkedin.com

Licurgo Benemann De Almeida
LinkedIn Corporation
New York, USA
lalmeida@linkedin.com

ABSTRACT

Computational marketing has become increasingly important in today's digital world, facing challenges such as massive heterogeneous data, multi-channel customer journeys, and limited marketing budgets. In this paper, we propose a general framework for marketing AI systems, the Neural Optimization with Adaptive Heuristics (NOAH) framework. NOAH is the first general framework for marketing optimization that considers both to-business (2B) and to-consumer (2C) products, as well as both owned and paid channels. We describe key modules of the NOAH framework, including prediction, optimization, and adaptive heuristics, providing examples for bidding and content optimization. We then detail the successful application of NOAH to LinkedIn's email marketing system, showcasing significant wins over the legacy ranking system. Additionally, we share details and insights that are broadly useful, particularly on: (i) addressing delayed feedback with lifetime value, (ii) performing large-scale linear programming with randomization, (iii) improving retrieval with audience expansion, (iv) reducing signal dilution in targeting tests, and (v) handling zero-inflated heavy-tail metrics in statistical testing.

CCS CONCEPTS

• **Computing methodologies** → **Machine learning**; • **Applied computing** → **Marketing**; • **Mathematics of computing** → *Probability and statistics; Mathematical optimization.*

*Corresponding author.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

KDD '24, August 25–29, 2024, Barcelona, Spain

© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 979-8-4007-0490-1/24/08

<https://doi.org/10.1145/3637528.3671591>

KEYWORDS

Computational Marketing, Linear Programming, Lifetime Value, Audience Expansion

ACM Reference Format:

Changshuai Wei, Benjamin Zelditch, Joyce Chen, Andre Assuncao Silva T Ribeiro, Jingyi Kenneth Tay, Borja Ocejo Elizondo, Sathiya Keerthi Selvaraj, Aman Gupta, and Licurgo Benemann De Almeida. 2024. Neural Optimization with Adaptive Heuristics for Intelligent Marketing System. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD '24)*, August 25–29, 2024, Barcelona, Spain. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3637528.3671591>

1 INTRODUCTION

Compared with traditional media marketing, online digital marketing has become increasingly important for advertisers and marketers. With online digital marketing, marketers can reach customers more directly with personalized marketing messages. Many of the current major internet companies build advertising platforms to better connect advertisers/marketers to customers, typically through auction or bidding systems. From a marketer's perspective, there are paid social channels (e.g., Facebook) and paid search channels (e.g., Google). Besides these paid channels, there are also owned channels like email, where marketers can connect with customers directly. There is much literature on building machine learning (ML) systems from ads platform's perspective; however, few papers have discussed ML and optimization systematically from the marketer's perspective. In fact, there are many unique challenges when building an intelligent marketing system from the marketer's perspective:

- (1) **The data can be massive and heterogeneous.** Since there are multiple marketing channels, the granularity of data from different channels are highly heterogeneous. For owned channels like email, marketers can collect member-level user feedback, whereas for paid channels like Google Search, marketers can only collect aggregate-level information, e.g., at the keyword or campaign level.

- (2) **The objective can be challenging to define.** The marketing life-cycle can be quite different across products. For many to-consumer (2C) products, the interaction is typically online and self-served, so the life-cycle from receiving the marketing message to conversion can be just a few days. For to-business (2B) products, the life-cycle can span weeks to months, including activities such as lead generation and sales involvement. As a result, these ML models face issues associated with delayed feedback.
- (3) **The action space is complex.** There are different types of actions to optimize the marketing outcome. These include (but are not limited to) at what time and frequency the marketing message should be sent, what content (e.g., picture, language) to use for the message, which audience to target, how much to bid for the marketing campaign(s) on various ads platforms. Unifying these actions under a general optimization framework would help build a more efficient marketing system.
- (4) **There are various types of constraints.** Marketing typically comes with cost. Some costs are explicit (e.g., billing from paid channels), while others are implicit (e.g., complaints or unsubscriptions from users on owned channels). Besides these cost constraints, there are also other operational constraints within marketing systems, e.g., capping on message frequency, minimum volume for various type of products. An intelligent marketing system should be able to determine optimal actions while respecting these constraints.

To address these challenges, we propose *Neural Optimization with Adaptive Heuristics (NOAH)*, a general framework for intelligent marketing systems. Before we introduce the framework, we briefly review related ML and optimization works in the marketing industry and for general recommender systems.

1.1 Related Work

1.1.1 ML in Marketing Industry. The use of ML in marketing has been fertile over the last decade, given the wealth of data naturally produced by marketing and the capacity of ML tools/methods to transform data into information that can be used for marketing optimization and decision-making. ML solutions have been proposed and/or applied to multiple domains, for instance: customer segmentation in industries such as retail [14, 32], hospitality [1, 44] and banking [33, 45]; customer lifetime value (LTV) [19] in industries ranging from gaming [11], to e-commerce [22], to subscription services [28]; marketing attribution, in both multi-touch attribution [38] and marketing mix modeling [23]; the optimization of different marketing channels, such as email [31, 40], paid search bidding [16, 18], display advertising [36, 41], among others. For a comprehensive review of the use of ML in marketing, we recommend the work presented in [26, 29].

1.1.2 Optimization in Recommender Systems. A commonly used optimization approach for Recommender systems (RecSys) is the ranking heuristic, where items are ranked by their relevance scores and top-ranked items are shown/sent to users. To balance different (and sometimes conflicting) objectives, the relevance score is often a linear combination of multiple predicted objective values. While a

large amount of research effort has gone into applying ML and AI to predicting the individual objective values, less work has been done on how to choose the weights for an optimized linear combination. In recent years, Linear Programming (LP) has been emerging in RecSys for its capability of trading off multiple objectives.

LP has been studied closely since the 1940s, and several commercial solvers such as Gurobi [20] and GLPK [30] can give high-quality solutions for small-to-moderate instances in reasonable time frames. However, they cannot scale to industrial-scale problems with billions or trillions of variables due to memory or computational constraints. One approach to handle large-scale LPs is to devise algorithms where certain key quantities can be computed in a distributed manner. The alternating direction method of multipliers (ADMM) [7] is one such approach; [46] details how ADMM can be implemented successfully for industrial-scale problems. Another approach is the primal-dual hybrid gradient (PDHG) method [9] for a class of problems in convex optimization. [2] applies the PDHG method to a saddle-point formulation of LP and adds enhancements such as diagonal preconditioning, presolving, adaptive step sizes and adaptive restarting to make the algorithm more efficient in practice.

1.2 Our Approach

NOAH is a general framework for intelligent marketing systems which consists of three major modules: 1) prediction module, 2) optimization module and 3) adaptive heuristics module. The prediction module provides predictions of business metrics given possible marketing actions. The optimization module chooses marketing actions that optimize the relevant business metrics. The adaptive heuristics module connects various components of the system to form a better feedback loop. We first go over more details of the NOAH framework in Section 2, including the three NOAH modules and some illustrative high-level examples. We then walk through a large scale application to LinkedIn’s marketing engine in Section 3. We also share production A/B test win results with some learnings in Section 4 and conclude in Section 5. We would like to highlight two major contributions of the paper:

- (1) As far as we know, this is the first general framework for marketing optimization that takes into account both 2B and 2C products, as well as both owned and paid channels.
- (2) This is a large-scale application to LinkedIn marketing that led to a major business win, and we report details of our learnings and experiences in this paper.

These contributions imply that our approach offers great scope for improving marketing optimization in general and hence is broadly useful.

2 NOAH FRAMEWORK

In this section, we will first give a brief overview of LinkedIn’s marketing system. We then introduce the details of different modules and components of the NOAH framework, followed by application guidelines and high-level examples for bidding and content optimization.

2.1 Marketing Ecosystem at LinkedIn

LinkedIn is the world’s largest professional network and has products that serve both 2C and 2B needs. As a result, the marketing system (Figure 1) needs to promote not only 2C products like Premium but also various 2B products in Talents, Learning, Ads and Sales lines of business. The marketing system is a heterogeneous system that involves both a first-party system and integration with third-party tools. Marketers interact with the marketing system to decide the optimal marketing actions for these 2B and 2C products, including whom the marketing message should be sent to, the best timing and frequencies for the messages to be sent, what bidding price or target the keywords or campaigns should be set at (on paid channels). Once the actions are decided from the Marketing Engine, the marketing messages are delivered via various owned or paid channels.

Paid channels include paid search, paid social and other display ads on third-party platforms, where LinkedIn participates as an advertiser in ad bidding through various auction mechanisms. Meanwhile, with owned channels, LinkedIn can promote its products on LinkedIn itself or use the email channel to have direct connections with customers. When a user clicks on a marketing message from one of these channels, they are directed to an optimized landing page where more detailed product information is shown. For 2C products, the landing page is typically on LinkedIn’s main website and the purchase can be typically completed in a self-serve manner. On the other hand, for 2B products the landing page is typically LinkedIn’s 2B micro-site (e.g., <https://business.linkedin.com/>) where an inquiry form can be generated for LinkedIn sales personnel to follow up.

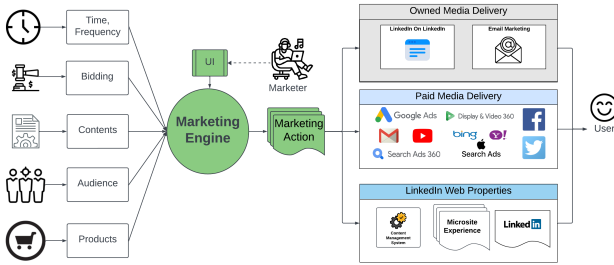


Figure 1: Overview of LinkedIn’s Marketing System.

2.2 NOAH Modules

In order to have a general framework that can handle different types of entities and possible actions in owned and paid channels, we first abstract out two definitions: *marketing units* and *marketing actions*. A *marketing unit* is an entity that a *marketing action* can be applied on. For example, in email marketing where we can directly connect with members and collect engagement data, the marketing unit is a member and a marketing action (to be optimized) can be whether to send a particular email to the member. In paid search channels, assuming marketers can get keyword-level data, the marketing unit can be a keyword, and a marketing action can be the bid amount for the keyword.

Figure 2 shows the overall diagram of the NOAH system. The core engine (inside the large dashed green box) consists of three key modules: (i) prediction, (ii) optimization, and (iii) adaptive heuristics. We first introduce the components outside the core engine so that the context, inputs and outputs around the core engine are clear.

NOAH is a human-in-the-loop system. Marketers with expert knowledge define appropriate marketing units and marketing actions in the relevant databases. A candidate retrieval process retrieves the appropriate set of marketing units and actions from the databases and feeds them to the core engine. Marketers also provide input parameters via an UI that can guide the core engine’s behavior. User engagement feedback and other contextual information are also collected as input to the core engine. The core engine processes these inputs through the three key modules to generate optimized marketing actions.

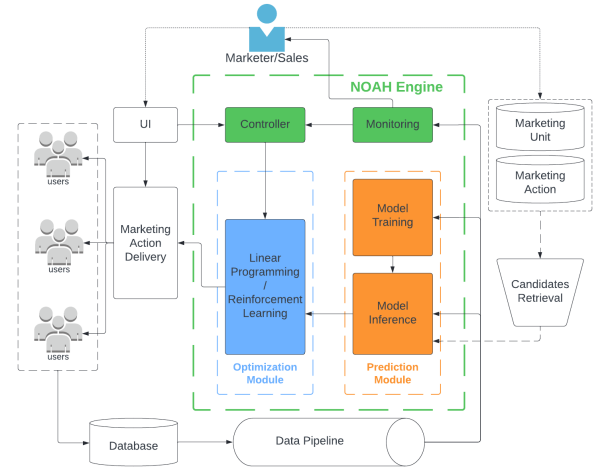


Figure 2: NOAH Framework.

2.2.1 Prediction Module. In the prediction module, we want to predict performance metrics for each marketing unit given some marketing action. Denoting the feature vector for marketing unit u as x_u , feature vector for marketing action a as x_a and a metric as y , the prediction module builds a functional mapping:

$$y = f(x_u, x_a, x_c),$$

where x_c represents a feature vector for contextual information other than marketing unit and action (e.g., location, device). The metrics are typically marketing objectives and constraints related to user feedback. For example, we can build a functional mapping for four metrics: click through rate (CTR), conversion rate, profit per conversion and cost per click (CPC). With these metrics, we can use the chain rule to connect these conditional probabilities to estimate the profit per impression and cost per impression, which can be further used in the optimization module. The functional mapping $f(\cdot)$ can be learnt from user feedback data. Depending on the amount of data available and the maturity of the system, we can use different supervised learning methods to estimate $f(\cdot)$. If the amount of data is small and the project is at the initial stage, a simple

regression model is often enough. If more data is available and the project is at the initial stage, an off-the-shelf gradient boosted tree [17] or XGBoost model [12] can give decent performance. If data is abundant and the project is in the iteration phase, we have found that deep neural networks give the best performance due to their flexible architecture. For example, transformer architectures can be used to generate embeddings for historical marketing interaction with users. If multiple metrics are being modeled, multi-task learning architectures like multi-gate mixture-of-experts (MMOE [25]) or entire space multi-task model (ESMM [27]) can be employed to exploit the relationship among multiple metrics during training.

One key characteristic of the prediction module is the inclusion of the marketing action. If the metric prediction is conditioned on the marketing action, the prediction module can then generate metric estimates for all possible marketing actions, from which the optimization module can choose the best action. However, historical data may not contain all possible marketing actions for each unit, which may generate bias in inference [10]. To reduce the bias, various causal ML methods can be used so that the $f(\cdot)$ learnt better reflects the causal relationship [15, 21, 39].

2.2.2 Optimization Module. With the functional mapping of marketing actions and performance metrics, we then need to choose the best action for a given marketing unit in the optimization module. The simplest, widely adopted approach is a ranking approach: choose the action with the largest y value, i.e., the largest estimated metric. While it can often work, this ranking approach ignores many limitations or constraints that can lead to sub-optimal and sometimes unacceptable performance.

To address various constraints across marketing units, we can frame the problem as a constrained optimization problem. For each metric $k = 0, 1, \dots, K$, let $y^k = f_k(x_u, x_{a_u}, x_c)$ denote the estimated value of metric k for marketing unit u under action a_u . Without loss of generality, let y^0 be the primary metric to be optimized, and let the remaining y^k ($k \geq 1$) be metrics that are subject to various limitations or business constraints with upper bound constraints C_k . Assuming no interaction effects across marketing units, we can formulate the optimization problem:

$$\max_{a_u} \sum_u f_0(x_u, x_{a_u}, x_c), \quad (1)$$

$$\text{s.t.} \sum_u f_k(x_u, x_{a_u}, x_c) \leq C_k, \quad \forall 1 \leq k \leq K. \quad (2)$$

The above optimization problem can be relaxed and solved with efficient linear programming (LP) algorithms in many cases; we detail some examples in Sections 2.3 and 3.3. When latency is a strong requirement, some compromises on accuracy need to be made, for example, to use yesterday's dual solution for today's primal calculation. Fortunately, for marketing AI use cases, latency is not a major problem as the systems are mostly offline. For example, in email marketing, marketers control the time and frequency and typically limit the frequency to avoid spamming users. In paid media, marketers only need to change the non-impression-level bid, e.g., cost-per-click (CPC) bid as the impression-level response is taken care of by the ads platform. For these reasons, LP is well suited for marketing AI use cases.

Besides the marketing cost and business constraints on the "space" dimensions, there are implicit constraints on the "time" dimension, particularly for each marketing unit. The time window to attract and convert a potential user is typically limited. Moreover, a marketing action taken at a given time point can impact the status and best marketing action to take at the next time point. We can leverage Reinforcement Learning (RL) to address optimization on "time" dimension. Contextual bandits can be used to balance the exploration and exploitation. Q-Learning can be used to help find an optimal sequence of marketing actions to maximize value for a given marketing unit.

To summarize, the optimization module addresses the resource limitation problem. As a general principle, we recommend LP when the resource limitation across marketing units is strict and there are less frequent marketing interactions, and recommend RL when there are more frequent marketing interactions and less constraints across marketing units (or these constraints can be taken care of by other parts of the system). We note that these two techniques are not mutually exclusive and can be combined. For example, we can focus on "space" constraints and use an LP formulation, while simplifying the time horizons resolution to an LTV prediction. We can also focus on "time" horizons and use a Q-learning formulation, while transforming the cost constraints implicitly to reward shaping. We argue that the former combination is often preferred in marketing, since marketing use cases are typically less frequent interaction with strict cost constraints.

2.2.3 Adaptive Heuristics Module. Adaptive heuristics (AH) are lightweight components that help the whole system run stably and coherently, each of which may have some prediction and optimization functionality. We here briefly discuss two main classes of AHs and provide more details in Appendix B.

One important class of AH is *Feedback Loop Controllers* for business constraints (green boxes as shown in Figure 2). Since the market is dynamic and the marketing AI system can have errors (e.g., prediction error), the observed cost C' may be greater (violation) or much less (under-utilization) than the budget C . To minimize under- or over-utilization, we can continuously monitor the performance and build a lightweight controller, $C^* = g(C, C')$, to update C^* as the input to the optimization module instead of the original C .

Another important class of AH is *Hierarchical Optimization*, so that the aggregate-level constraints propagate to lower level. For example, a macro-level optimization with a forecasting model and an LP can be built to allocate marketing budget to different marketing channels and products.

2.3 Guidelines and Examples

NOAH is a general framework for marketing AI systems. Ideally, we would like to perform joint optimization for multiple types of marketing actions, e.g., product, timing, frequencies and content. However, this is often infeasible due to reasons such as infrastructure constraints, operational convenience, or limitations in the ML models or optimization solvers. As a workaround, we can use the same NOAH framework and perform optimization in a cascading/sequential manner for different types of marketing actions. For example, we can first decide the product for each marketing unit, then optimize the timing and frequencies, and finally choose the

best content. Another benefit of this multi-stage design is that it allows for a flexible combination of different optimization techniques for different actions. Here, we give examples for two different marketing actions: bidding and content optimization.

2.3.1 Bidding Optimization Example. For bidding, advertisers can receive hourly or daily ads winning performance metrics x_p , for example, average position for each keyword, or impression winning percentage for each campaign. One can build a series of models to predict: number of clicks $y^{\text{clicks}} = f_{\text{clicks}}(x_u, x_p)$, cost per click $y^{\text{CPC}} = f_{\text{CPC}}(x_u, x_p)$, conversion rate $y^{\text{CVR}} = f_{\text{CVR}}(x_u)$ and revenue per conversion $y^{\text{RPC}} = f_{\text{RPC}}(x_u)$. Note that we can assume y^{CVR} and y^{RPC} do not depend on ads winning performance as they happen after users click the ad and land on the landing page. Typically, advertisers want to maximize revenue while maintaining a certain level of return-on-ad-spend (ROAS). The corresponding constrained optimization formulation could be:

$$\begin{aligned} \max_{p(u)} \quad & \sum_u y_{u,p(u)}^{\text{rev}}, \\ \text{s.t.} \quad & \frac{\sum_u y_{u,p(u)}^{\text{rev}}}{\sum_u y_{u,p(u)}^{\text{cost}}} \geq C_{\text{ROAS}}, \end{aligned}$$

where p represents the index for categorized winning performance, and

$$\begin{aligned} y_{u,p}^{\text{rev}} &= f_{\text{clicks}}(x_u, x_p) f_{\text{CVR}}(x_u) f_{\text{RPC}}(x_u), \\ y_{u,p}^{\text{cost}} &= f_{\text{clicks}}(x_u, x_p) f_{\text{CPC}}(x_u, x_p). \end{aligned}$$

The constrained optimization problem can be solved with the efficient algorithm in Appendix A.1 if the scale is large. Once the desired optimal x_p^* is obtained, we can calculate the optimal CPC bid as $a_u^* = f_{\text{CPC}}(x_u, x_p^*)$, which can be directly used in first-price auction platforms, or as the expected-cost-per-click (ECPC) bid in second-price auction platforms.

2.3.2 Content Optimization Example. Content optimization can happen at different stages across different marketing channels. Without loss of generality, assume there is a slot on the landing page where we can show different creatives a to different users u . We can build an online high-dimensional contextual bandit model based on following regression model:

$$\rho^{-1}(y) = x^T w + \epsilon, \quad \epsilon \sim N(0, \beta^2),$$

where $x = (x_u, x_a, x_{u,a})$ is feature vector that includes features for user x_u , features for creative x_a , and features for the user-creative interaction $x_{u,a}$, $\rho(\cdot)$ is a link function, and y is the observed reward. Online updates for the posterior distribution of $w_i \sim N(\mu_{w_i}, \sigma_{w_i}^2)$ can be computed efficiently using factor graph and expectation propagation (see algorithm details in Appendix A.2). Thompson sampling can then be employed to obtain predicted rewards for different creatives, and the user is shown the creative with the largest reward on the landing pages.

3 EMAIL MARKETING APPLICATION

In this section, we report the details of a large-scale application of the NOAH framework to LinkedIn's email marketing engine, including important innovations and learnings.

3.1 Email System

We can think of the email marketing system as a two-stage recommender system: candidate retrieval followed by email optimization (Figure 5 in Appendix). In the candidate retrieval stage, marketers create email marketing campaigns for different products and define the audience list for each campaign. In the email optimization stage, ML models are used to find the best marketing emails for a selected group of members, so as to optimize metrics while minimizing negative member experiences like spamming. The whole system runs in batches, with emails sent out at regular time intervals (e.g., weekly). The system seeks to maximize value for LinkedIn while respecting multiple business constraints. Some constraints are global, e.g., the total number of user unsubscriptions should be less than a certain value each week, while other constraints are local, e.g., each member cannot receive more than certain number of emails each week.

3.2 Prediction with LTV

LinkedIn faces distinct challenges in predicting marketing performance due to the different nature of 2C and 2B segments, particularly the extended decision-making processes and delayed feedback for 2B products. To overcome this, we create separate models for predicting short-term metrics and LTV, then integrate the results to estimate the long-term metric. Here, the long-term metric q^t is a numeric measure defined by marketers reflecting marketing performance over an extended period t , with short-term metrics representing observable metrics within a brief timeframe, like conversion and unsubsubscription. Notably, conversion is typically defined differently for 2B and 2C products. In 2B, it often refers to the submission of a qualified lead, whereas in 2C, a conversion can be marked by a customer signing up for a free trial.

To predict the short-term metrics, we develop predictive models forming a functional mapping:

$$\begin{aligned} y_{u,a}^{\text{conv}} &= f_{\text{conv}}(x_u, x_a, x_c), \\ y_{u,a}^{\text{unsub}} &= f_{\text{unsub}}(x_u, x_a, x_c), \end{aligned}$$

where y^{conv} and y^{unsub} represent the predicted conversion probabilities and unsubsubscription probabilities respectively, x_u denotes the feature vectors of users capturing user profiles, demographics, and other behavioral data, x_a symbolizes the marketing actions undertaken, including the decision to send specific marketing emails to certain members, and x_c represents other relevant contextual information that might influence the metrics.

We have a dedicated model to predict the LTV associated with a given conversion event. We calculate the LTVs as

$$y_{u,p}^{\text{lrv}} = f_{\text{lrv}}(x_u, x_p, x_c)$$

where p indexes the product that email campaign a tries to promote. We model the real-valued y^{lrv} using regression equipped with a gamma loss function, defined as

$$L = \sum_i \left[\mu_i k + q_i^{12} \exp(-\mu_i) k \right]$$

where k is the shape parameter, i indexes the observations in the training data, q_i^{12} is the observed 12-month long-term metric (the ground truth value) and $\mu_i = \log(y_i^{\text{lrv}})$. This is a long time window

that makes recent positive engagement unavailable for training (i.e., censored observations). In order to overcome this issue, we make auxiliary LTV predictions at shorter time windows (e.g., 1,3,6 months) and use the output of those models as features for the main 12-month LTV prediction. We found that this strategy increases prediction accuracy and allows the model to respond faster to evolving business trends.

3.3 Constrained and Multi-Objective Optimization

3.3.1 Formulating the Constrained Optimization Problem. In our legacy system, a ranking heuristic was used in the email optimization stage. For each member, all candidate emails were ranked by a linear combination of y^{conv} and y^{unsub} with fixed coefficient κ , i.e., $y^{\text{conv}} - \kappa y^{\text{unsub}}$, and top-ranked emails were sent to the member with a fixed frequency cap. This approach did not optimize the business goal directly, and business constraints were only considered implicitly.

Here we apply the constrained optimization formulation in Section 2.2.2 to this problem. Let $u = 1, \dots, U$ and $j = 1, \dots, J$ index the member and campaigns respectively, and let $a_{u,j} \in \{0, 1\}$ denote the binary action decision of whether to send email j to member u . We can estimate the total expected long-term metric by $\sum_{u,j} a_{u,j} y_{u,j}^{\text{conv}} y_{u,p(j)}^{\text{ltv}}$, and the total expected unsubscriptions by $\sum_{u,j} a_{u,j} y_{u,j}^{\text{unsub}}$.

Assuming that (i) we allow a maximum of C_{fcap} emails to be sent to each member, (ii) we must send at least C_{2B} and C_{2C} emails for 2B and 2C products respectively, and (iii) we must have no more than C_{unsub} unsubscriptions, we can form the constrained optimization problem as follows:

$$\begin{aligned} \min_{a_{u,j}} & \sum_{u,j} -y_{u,j}^{\text{conv}} y_{u,p(j)}^{\text{ltv}} a_{u,j}, \\ \text{s.t.} & \sum_{u,j} a_{u,j} y_{u,j}^{\text{unsub}} \leq C_{\text{unsub}}, \\ & \sum_u \sum_{j \in \mathbb{J}_{2B}} -a_{u,j} \leq -C_{2B}, \\ & \sum_u \sum_{j \in \mathbb{J}_{2C}} -a_{u,j} \leq -C_{2C}, \\ & \sum_j a_{u,j} \leq C_{\text{fcap}}, \quad \forall u, \\ & a_{u,j} \in \{0, 1\}, \quad \forall u, j. \end{aligned}$$

where, \mathbb{J}_{2B} and \mathbb{J}_{2C} denote the email campaigns for 2B and 2C products respectively.

We recognize the optimization problem as an integer linear program (ILP), which can be combinatorially difficult to solve [37]. Instead, we relax the constraints on the $a_{u,j}$'s to allow them to take on values in the interval $[0, 1]$. The resulting problem is a linear program (LP) which is tractable (details in Section 3.3.2). We can interpret non-integer $a_{u,j}$ as the probability with which we should send email j to member u .

3.3.2 Solving the Large-Scale LP. While there are open-source and commercial solvers for solving LPs, the size of the LP for our application is too large for these solvers. Recall that the optimization

variable a has dimension UJ , the product of the number of members and the number of campaigns. For a large social network like LinkedIn, we could have 10-100 millions of members and hundreds of email campaigns, resulting in tens of billions of variables. To solve the LP, we turn to *DualLP* [4, 34, 35], a large-scale LP solver that we developed at LinkedIn. We briefly outline the algorithm here.

The LP in 3.3.1 can be written in the following form:

$$\min_a y^\top a \quad \text{s.t.} \quad Da \leq b, \quad a_u \in \mathcal{B}_u, \quad u \in \{1, \dots, U\}, \quad (3)$$

where $a = (a_{1,1}, \dots, a_{1,J}, \dots, a_{U,1}, \dots, a_{U,J}) \in \mathbb{R}^{UJ}$ is the vector of decision variables, y is the coefficient vector for the decision variables in the objective, D is the design matrix in the constraints, b is the constraint budget vector, and \mathcal{B}_u are additional "simple" constraints, i.e., it is efficient to compute the Euclidean projection onto \mathcal{B}_u . In particular, $b = (C_{\text{unsub}}, -C_{2B}, -C_{2C})$, and Da corresponds to the left-hand side of the first three constraints from the previous section. \mathcal{B}_u corresponds to the remaining two constraints, i.e., box-cut constraints $\mathcal{B}_u = \{a_{u,j} : 0 \leq a_{u,j} \leq 1 \forall j, \sum_j a_{u,j} \leq C_{\text{fcap}}\}$.

Instead of solving the LP (3) directly, we solve a perturbed version of the LP, which is a quadratic program (QP):

$$\min_a y^\top a + \frac{\gamma}{2} a^\top a \quad \text{s.t.} \quad Da \leq b, \quad a_u \in \mathcal{B}_u, \quad u \in \{1, \dots, U\}, \quad (4)$$

where $\gamma > 0$ is a regularization hyperparameter. Dualizing just the polyhedral constraint $Da \leq b$ leads to the partial Lagrangian dual

$$g_\gamma(\lambda) = \min_{a \in \mathcal{B}} \left\{ y^\top a + \frac{\gamma}{2} a^\top a + \lambda^\top (Da - b) \right\},$$

where $\mathcal{B} = \prod_{u=1}^U \mathcal{B}_u$. The problem is transformed to calculate $\max_\lambda g_\gamma(\lambda)$ such that $\lambda \geq 0$, assuming strong duality holds. Applying Danskin's theorem [6], we know that g_γ is differentiable and can explicitly compute its derivative:

$$\begin{aligned} \nabla g_\gamma(\lambda) &= D\hat{a}(\lambda) - b, \\ \hat{a}_u(\lambda) &= \Pi_{\mathcal{B}_u} \left[-\frac{1}{\gamma} (D^\top \lambda + y)_u \right], \end{aligned} \quad (5)$$

where $\Pi_{\mathcal{B}_u}(\cdot)$ is the Euclidean projection operator onto \mathcal{B}_u , and for a vector z with the same dimensions as decision variable a , z_u denotes the sub-vector with indices corresponding to member u . Notice that the computation of $\hat{a}(\lambda)$ can be parallelized across members u and the box-cut projection $\Pi_{\mathcal{B}_u}$ can be computed efficiently by leveraging Wolfe's algorithm [43]. As a result, $\nabla g_\gamma(\lambda)$ can be computed efficiently. This allows us to use first-order methods such as accelerated gradient ascent [5] and L-BFGS [24] to obtain λ_γ^* , the dual value that optimizes the dual function g_γ . From this optimal dual value, we can obtain optimal primal values a_γ^* via (5):

$$(a_\gamma^*)_u = \hat{a}_u(\lambda_\gamma^*) = \Pi_{\mathcal{B}_u} \left[-\frac{1}{\gamma} (D^\top \lambda_\gamma^* + y)_u \right], \quad u = 1, \dots, U.$$

While we obtain the optimal solution for the perturbed problem and not the original one, [4] shows that for sufficiently small γ values, the perturbed solution is also optimal for the original problem.

3.3.3 Practical Tips for Implementation. We describe a couple of tips that make DuaLip converge much faster. (A) The choice of γ is associated with a critical trade-off: DuaLip’s solving of (4) is faster if γ is larger, whereas DuaLip’s approximate solution of (4) in a given number of iterations is closer to the solution of (3) if γ is smaller. For one run, we try a grid of γ values and choose the largest γ value such that the regularizer part of the objective is small enough, i.e., the ratio, $\frac{\gamma a^\top a}{2\|y^\top a\|}$ is smaller than some threshold, say 10^{-3} . After finding such a γ value, we fix it and use it for all future production runs. (B) If D is ill-conditioned (which typically happens when different constraints are on different scales), we can re-scale the constraints, i.e., pre-multiply D and b by a suitably chosen diagonal matrix.

With Dualip, the primal solution will fall on a vertex of \mathcal{B}_u with high probability [35] when γ is sufficiently small. This means most of the $(a_\gamma^*)_u$ ’s are already binary, i.e., in $\{0, 1\}$. However, we still need to handle the cases when they are between 0 and 1. Simple rounding may violate the hard constraint C_{cap} . We devise a randomization algorithm, basically sampling the campaign indexes with probability calculated from the primal solution (Appendix A.3). The randomization also provides the additional benefit of balancing exploration and exploitation.

3.4 Audience Expansion

In the email marketing lifecycle, dedicated marketing teams from each business unit provide a list of candidate members for each email campaign. Historically, we have found these candidate member lists conservative in scope. To address this, we augment the provided lists with additional members through an Audience Expansion (AE) module in which we retrieve additional members who resemble the members provided by the marketing teams and append them to the candidate lists. We find such members via an approximate nearest neighbors (ANN) algorithm.

Formally, let \mathcal{U} denote the overall member base at LinkedIn, let \mathcal{J} denote the set of marketing campaigns, and let $\mathcal{U}_j^o \subset \mathcal{U}$ denote the set of original candidate members provided by a marketing team for email marketing campaign $j \in \mathcal{J}$. We leverage ANN via locality-sensitive hashing [13] to produce a set

$$\mathcal{U}_j^{\text{expansion}} \subseteq \left\{ u \in \mathcal{U}_j^- : \exists u' \in \mathcal{U}_j^o, \text{ s.t. } \|x_u - x_{u'}\|_2 \leq \delta \right\},$$

where, $\mathcal{U}_j^- = \mathcal{U} \setminus \mathcal{U}_j^o$ and δ is a hyperparameter which controls the size of the expansion audience. Ultimately, we combine the original audience with expansion audience for each campaign and input the expanded audience \mathcal{U}^e into our optimization engine, where $\mathcal{U}^e = \bigcup_{j \in \mathcal{J}} \mathcal{U}_j^o \cup \mathcal{U}_j^{\text{expansion}}$.

We have found that as we relax δ , there is a trade-off between wider campaign reach through increased expansion audience size and diminished incremental expansion audience quality. If δ is too large, a portion of the incremental expansion members added through AE are so dissimilar to the campaign’s original members that they are never chosen by the LP to receive the marketing campaign. Moreover, the computational runtime of both Audience Expansion and the LP optimization increase in δ , so choosing δ too large can lead to additional runtime for minimal gain. This latter concern about computational efficiency is especially important

at LinkedIn’s scale. On the other hand, if δ is too small then the expansion audience itself will be too small to produce a significant gain in downstream metrics of interest. We tune δ through offline evaluation to balance these trade-offs.

4 EXPERIMENTS AND RESULTS

4.1 Offline Analysis

We performed extensive offline analysis of prediction models, LP and AE. Due to space limitation, we only highlight the summary of offline analysis here and share details of the results in Appendix D of the arXiv version of the paper [42].

- **Evaluation of prediction models.** We evaluated the classification and calibration performance of the two propensity models f_{conv} (AUC=0.895) and f_{unsub} (AUC=0.715). We also evaluated the LTV model f_{lrv} and found that: (i) gamma loss with auxiliary tasks can respond faster to market changes, (ii) it performs better than an alternative modeling approach using survival loss.
- **LP analysis.** We found that smaller γ values lead to better optimality but can take longer to converge. We also validated the high probability of binary primal solutions. More importantly, we investigated the fallback strategy of using the previous dual solution for the current LP, and found that we can maintain 99% of the objective value if data distribution is stable.
- **AE analysis.** We evaluated the impact of δ with respect to the trade-off between audience size, quality and runtime.
- **Pre-A/B test simulation.** We find that (i) NOAH can make efficient “not-send” decisions through LP whereas legacy ranking system cannot, and (ii) the combination of AE and LP helps realize the full potential of the optimization engine.

4.2 Online Test Design

We designed an online A/B test to compare the performance of the NOAH email system with that of the legacy system. In the test, we randomly assigned members to the treatment arm (NOAH system) or the control arm (legacy system). In addition, we included a third experimentation arm for the legacy system equipped with AE to isolate the impact of increasing marketing email volume on its own. The flow of the test is visualized in Figure 3.




Figure 3: A/B Test Flow.

The email marketing channel only covers a small percentage of the member base. If the randomization splits on the set of all members, the treatment effect could be diluted. We therefore focus on members that are “available” to receive marketing emails. In particular, we first randomize all members to 3 buckets, $\mathcal{U}_b^{\text{split}}$,

where $b \in \{1, 2, 3\}$. For each week t of the experiment, we can expand the original audience \mathcal{U}_t^o to \mathcal{U}_t^e with AE. Assuming the experiment runs for T weeks, the set of members for experiment arm b can be obtained as $\mathcal{U}_{b,T} = \bigcup_{t \leq T} \mathcal{U}_b^{\text{split}} \cap \mathcal{U}_t^e$. Note here that for control arm $b = 1$, the legacy system will rank and send emails only for members in original audience, i.e., over set of $\mathcal{U}_t^o \cap \mathcal{U}_1^{\text{split}}$ for week t .

Since it is not realistic to wait for one year to obtain the long-term metric q^{12} , we define a long-term surrogate index [3] $\mathcal{Y}_u = \sum_{p \in \mathcal{P}_u} y_{u,p}^{\text{LTV}}$ for each member u , where \mathcal{P}_u denote the set of products that member u converted during the test period, and $y_{u,p}^{\text{LTV}}$ denotes the expected 1-year long-term metric associated with the conversion on product p for member u . It is worth noting that although each observation consists of potentially multiple conversions, through offline analysis we have discovered that the majority of members convert at most one time during the test period.

Besides the above primary metric, we also monitor a few secondary metrics, e.g., email volume, conversions (Conv) as well as observed long-term metric. In addition, we have other metrics corresponding to the LP constraints, the key one being unsubscriptions (Unsub), a binary metric indicating whether or not a user unsubscribed during the test period.

We compare the means of the above metrics from different arms for statistical testing of treatment effects. It is worth noting that long-term metrics exhibit characteristics of a zero-inflated heavy-tail distribution (Appendix C.1). We performed simulations on Type I error and power of the t -test against these distributions, and found that the standard t -test is actually robust (Appendix C.2). We believe this result is broadly useful, as alternatives do not suit the needs of such testing: (1) Winsorization removes the large data points that actually matter (if not more), (2) Non-parametric tests (e.g., Mann-Whitney U test) lack a corresponding treatment effect estimate.

4.3 Online Test Results

4.3.1 Overview. The A/B test resulted in major wins, one of the largest in LinkedIn algorithmic marketing thus far. An overview of the results is presented in Table 1. All the numbers in the tables represent **relative** difference over the control arm, i.e., ratio of the difference (between the treatment and control arms) to the mean value in control. Results that are not statistically significant are marked either "neutral" or "not stat sig".

Table 1: Overview of A/B Test Result

Marketing System	Surrogate Index	Unsub	Email Volume
Legacy+AE	neutral	+8.08%	+9.42%
NOAH	+2.16% ¹	+1.52%	+8.44%

¹ All percentages are relative to the control arm.

As shown in Table 1, NOAH produces a statistically significant positive lift in expected 1-year long-term metric (Surrogate Index) at the cost of a moderate uptick (+1.52%) in total unsubsubscription. Given that Email Volume is increased by 8.44%, the unsubsubscription per email sent is actually reduced, indicating that more relevant

Table 2: Product Type Breakdown

Product Type	Surrogate Index	3-Mth Metric ¹	Conv
2C	not stat sig	neutral	neutral
2B	+13.21%	+23.66%	+13.30%

¹ Observed 3-month long-term metric q^3 after test concluded.

emails were sent. By comparison, the legacy ranking system with AE (Legacy+AE) does not produce a statistically significant effect on the long-term metric, and it produces a large Unsub increment (+8.08%), a magnitude similar as Email Volume increment (+9.42%).

4.3.2 Product Type Breakdown. Historically, we have faced challenges in developing a unified marketing system which optimizes marketing emails for a heterogeneous suite of 2C and 2B products. Our legacy marketing system tends to optimize for quick sign-up 2C products in spite of these products' relatively small lifetime values. On the other hand, campaigns associated with high-LTV 2B products tend to be neglected.

A breakdown of comparison between NOAH and control into 2C and 2B product types is presented in Table 2. We observe that NOAH is able to drive a large positive lift in both Surrogate Index (+13.21%) and total Conv (+13.30%) for high-LTV 2B products without sacrificing these metrics for 2C products. It achieves this lift by increasing total 2B conversions and shifting the distribution 2B conversions toward those with higher Surrogate Index, as visualized in Figure 7 in the Appendix.

Due to the nature of delayed feedback, we tracked the long-term metric of conversions that happened during A/B test even after the test concluded. We find that NOAH is able to drive a significant lift (+23.66%) in observed 3-month long-term metric (3-Mth Metric) for 2B products. As with the Surrogate Index, we observed a distributional shift toward high-value conversions in 3-Mth Metrics, as shown in Figure 8 in the Appendix.

The breakdown analysis demonstrates that NOAH is able to lift high value long-term metrics for 2B products while simultaneously maintaining 2C acquisition metrics and controlling unsubscriptions. This holistic behavior is essential for a unified marketing system which must support a wide range of campaigns having different downstream values.

5 CONCLUSION

In this paper, we introduce various challenges faced when building an intelligent marketing system and propose a general framework, NOAH, to address these challenges. We give guidelines on how to apply the NOAH framework and provide a few high-level examples. In addition, we report a large-scale application of NOAH to LinkedIn's email marketing use case. We discuss important innovations and learnings in the implementation and share results of the major wins for the corresponding online A/B test.

ACKNOWLEDGMENTS

We thank Shruti Sharma, Maggie Zhang, Divyakumar Menghani, Aarthi Jayaram, and Rehan Khan for their support of the work. We also thank our colleagues Michael Ingley, Grinishkumar Engineer,

Artem Baranov, Sumukh Sagar Manjunath, Carlos Flores, Shawn Bianchi, Kim Foo, Kyla Falkenhagen for their collaboration on the email application and A/B test. Additionally, we'd like to acknowledge Zhiqi Guo, Ming Wu, Miao Cheng and Faisal Farooq for their contributions and support during the early phase of the work.

REFERENCES

- [1] Elaheh Yadegaridehkordi Louis Sanzogni A. Rashid Tarik Kathy Knox Sarminah Samad Othman Ibrahim Ali Ahani, Mehrbakhsh Nilashi. 2002. Revealing customers' satisfaction and preferences through online review analysis: The case of Canary Islands hotels. *Journal of Retailing and Consumer Services* 51, 3 (2002), 331–343.
- [2] David Applegate, Mateo Diaz, Oliver Hinder, Haihao Lu, Miles Lubin, Brendan O'Donoghue, and Warren Schudy. 2021. Practical Large-Scale Linear Programming using Primal-Dual Hybrid Gradient. 20243–20257.
- [3] Susan Athey, Raj Chetty, Guido W Imbens, and Hyunseung Kang. 2019. *The surrogate index: Combining short-term proxies to estimate long-term treatment effects more rapidly and precisely*. Technical Report. National Bureau of Economic Research.
- [4] Kinjal Basu, Amol Ghoting, Rahul Mazumder, and Yao Pan. 2020. ECLIPSE: An Extreme-Scale Linear Program Solver for Web-Applications. *Proceedings of the 37th International Conference on Machine Learning*, 704–714.
- [5] Amir Beck and Marc Teboulle. 2009. A Fast Iterative Shrinkage-Thresholding Algorithm for Linear Inverse Problems. *SIAM Journal on Imaging Sciences* 2 (2009), 183–202. Issue 1. <https://doi.org/10.1137/080716542>
- [6] Dimitri P Bertsekas. 1997. Nonlinear programming. *Journal of the Operational Research Society* 48, 3 (1997), 334–334.
- [7] Stephen Boyd, Neal Parikh, Eric Chu, Borja Peleato, and Jonathan Eckstein. 2011. Distributed optimization and statistical learning via the alternating direction method of multipliers. *Foundations and Trends in Machine Learning* 3 (2011), 1–122. Issue 1. <https://doi.org/10.1561/22000000016>
- [8] Eric Breck, Shanjing Cai, Eric Nielsen, Michael Salib, and D Sculley. 2017. The ML test score: A rubric for ML production readiness and technical debt reduction. In *2017 IEEE International Conference on Big Data (Big Data)*. IEEE, 1123–1132.
- [9] Antonin Chambolle and Thomas Pock. 2011. A first-order primal-dual algorithm for convex problems with applications to imaging. *Journal of Mathematical Imaging and Vision* 40 (5 2011), 120–145. Issue 1. <https://doi.org/10.1007/s10851-010-0251-1>
- [10] Jiawei Chen, Hande Dong, Xiang Wang, Fuli Feng, Meng Wang, and Xiangnan He. 2023. Bias and debias in recommender system: A survey and future directions. *ACM Transactions on Information Systems* 41, 3 (2023), 1–39.
- [11] Pei Pei Chen, Anna Guitart, Ana Fernández del Río, and Africa Perianez. 2018. Customer lifetime value in video games using deep learning and parametric models. In *2018 IEEE international conference on big data (big data)*. IEEE, 2134–2140.
- [12] Tianqi Chen and Carlos Guestrin. 2016. Xgboost: A scalable tree boosting system. In *Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining*. 785–794.
- [13] Mayur Datar, Nicole Immorlica, Piotr Indyk, and Vahab S. Mirrokni. 2004. Locality-sensitive hashing scheme based on p-stable distributions. *Proceedings of the twentieth annual symposium on Computational geometry* (2004), 253–262.
- [14] Michelle Roehm Derrick S. Boone. 2002. Retail segmentation using artificial neural networks. *International Journal of Research in Marketing* 19, 3 (2002), 287–301.
- [15] Max H Farrell, Tengyuan Liang, and Sanjog Misra. 2020. Deep learning for individual heterogeneity: An automatic inference framework. *arXiv preprint arXiv:2010.14694* (2020).
- [16] Zhe Feng, Swati Padmanabhan, and Di Wang. 2023. Online Bidding Algorithms for Return-on-Spend Constrained Advertisers. In *Proceedings of the ACM Web Conference 2023*. 3550–3560.
- [17] Jerome H Friedman. 2001. Greedy function approximation: a gradient boosting machine. *Annals of statistics* (2001), 1189–1232.
- [18] Hui Gao and Yihan Yang. 2022. Multi-Head Online Learning for Delayed Feedback Modeling. *arXiv preprint arXiv:2205.12406* (2022).
- [19] Sunil Gupta, Dominique Hanssens, Bruce Hardie, Wiliam Kahn, V Kumar, Nathaniel Lin, Nalini Ravishanker, and S Sriram. 2006. Modeling customer lifetime value. *Journal of service research* 9, 2 (2006), 139–155.
- [20] Gurobi Optimization, LLC. 2023. Gurobi Optimizer Reference Manual. <https://www.gurobi.com>
- [21] Daniel Jacob. 2021. Cate meets ml: Conditional average treatment effect and machine learning. *Digital Finance* 3, 2 (2021), 99–148.
- [22] Lukáš Kakalejčík, Jozef Bucko, and Martin Vejačka. 2019. Differences in buyer journey between high- and low-value customers of e-commerce business. *Journal of theoretical and applied electronic commerce research* 14, 2 (2019), 47–58.
- [23] Ajay Kumar, Ravi Shankar, and Naif Radi Aljohani. 2020. A big data driven framework for demand-driven forecasting with effects of marketing-mix variables. *Industrial marketing management* 90 (2020), 493–507.
- [24] Dong C. Liu and Jorge Nocedal. 1989. On the Limited Memory BFGS Method for Large Scale Optimization. *Mathematical Programming* 45 (1989), 503–528.
- [25] Jiaqi Ma, Zhe Zhao, Xinyang Yi, Jilin Chen, Lichan Hong, and Ed H Chi. 2018. Modeling task relationships in multi-task learning with multi-gate mixture-of-experts. In *Proceedings of the 24th ACM SIGKDD international conference on knowledge discovery & data mining*. 1930–1939.
- [26] Liye Ma and Baohong Sun. 2020. Machine learning and AI in marketing—Connecting computing power to human insights. *International Journal of Research in Marketing* 37, 3 (2020), 481–504.
- [27] Xiao Ma, Liqin Zhao, Guan Huang, Zhi Wang, Zelin Hu, Xiaoqiang Zhu, and Kun Gai. 2018. Entire space multi-task model: An effective approach for estimating post-click conversion rate. In *The 41st International ACM SIGIR Conference on Research & Development in Information Retrieval*. 1137–1140.
- [28] Edward C Malthouse and Robert C Blattberg. 2005. Can we predict customer lifetime value? *Journal of interactive marketing* 19, 1 (2005), 2–16.
- [29] Eric WT Ngai and Yuanyuan Wu. 2022. Machine learning in marketing: A literature review, conceptual framework, and research agenda. *Journal of Business Research* 145 (2022), 35–48.
- [30] Eiji Oki. 2012. GLPK (GNU Linear Programming Kit). <https://api.semanticscholar.org/CorpusID:63578694>
- [31] M Paulo, Vera L Miguéis, and Ivo Pereira. 2022. Leveraging email marketing: Using the subject line to anticipate the open rate. *Expert systems with applications* 207 (2022), 117974.
- [32] Musthofa Galih Pradana and Hoang Thi Ha. 2021. Maximizing strategy improvement in mall customer segmentation using k-means clustering. *Journal of Applied Data Sciences* 2, 1 (2021), 19–25.
- [33] Omri Raiter. 2021. Segmentation of bank consumers for artificial intelligence marketing. *International Journal of Contemporary Financial Issues* 1, 1 (2021), 39–54.
- [34] Rohan Ramanath, Sathya S. Keerthi, Kinjal Basu, Konstantin Salomatin, Pan Yao, Amol Ghoting, and Miao Cheng. 2022. DualLip: Dual Decomposition based Linear Program Solver, version 2.0.0. <https://github.com/linkedin/dualip>
- [35] Rohan Ramanath, S. Sathya Keerthi, Yao Pan, Konstantin Salomatin, and Kinjal Basu. 2021. Efficient Vertex-Oriented Polytopic Projection for Web-scale Applications. *arXiv preprint arXiv:2103.05277*, 1–13. <http://arxiv.org/abs/2103.05277>
- [36] Kan Ren, Weinan Zhang, Ke Chang, Yifei Rong, Yong Yu, and Jun Wang. 2017. Bidding machine: Learning to bid for directly optimizing profits in display advertising. *IEEE Transactions on Knowledge and Data Engineering* 30, 4 (2017), 645–659.
- [37] A. Schrijver. 1986. *Theory of Linear and Integer programming*. Wiley-Interscience.
- [38] Xuhui Shao and Lexin Li. 2011. Data-driven multi-touch attribution models. In *Proceedings of the 17th ACM SIGKDD international conference on Knowledge discovery and data mining*. 258–264.
- [39] Claudia Shi, David Blei, and Victor Veitch. 2019. Adapting neural networks for the estimation of treatment effects. *Advances in neural information processing systems* 32 (2019).
- [40] Duncan Simester, Artem Timoshenko, and Spyros I Zoumpoulis. 2020. Targeting prospective customers: Robustness of machine-learning methods to typical data challenges. *Management Science* 66, 6 (2020), 2495–2522.
- [41] Vinay Singh, Brijesh Nanavati, Arpan Kumar Kar, and Agam Gupta. 2023. How to maximize clicks for display advertisement in digital marketing? A reinforcement learning approach. *Information Systems Frontiers* 25, 4 (2023), 1621–1638.
- [42] Changshuai Wei, Benjamin Zelditch, Joyce Chen, Andre Assuncao Silva T Ribeiro, Jingyi Kenneth Tay, Borja Ochoa Elizondo, Keerthi Selvaraj, Aman Gupta, and Licurgo Benemann De Almeida. 2024. Neural Optimization with Adaptive Heuristics for Intelligent Marketing System. *arXiv:2405.10490*
- [43] Philip Wolfe. 1976. Finding the nearest point in a polytope. *Mathematical Programming* 11 (1976), 128–149.
- [44] Elaheh Yadegaridehkordi, Mehrbakhsh Nilashi, Mohd Hairul Nizam Bin Md Nasir, Saadeh Momtazi, Sarminah Samad, Eko Supriyanto, and Fahad Ghabban. 2021. Customers segmentation in eco-friendly hotels using multi-criteria and machine learning techniques. *Technology in Society* 65 (2021), 101528.
- [45] Danuta Zakrzewska and Jan Murewski. 2005. Clustering algorithms for bank customer segmentation. In *5th International Conference on Intelligent Systems Design and Applications (ISDA'05)*. IEEE, 197–202.
- [46] Jun Zhou, Yang Bao, Daohong Jian, and Hua Wu. 2023. PDAS: A Practical Distributed ADMM System for Large-Scale Linear Programming Problems at Alipay. *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 5717–5727. <https://doi.org/10.1145/3580305.3599883>

A ALGORITHMS

A.1 Bidding Optimization with Lagrangian Relaxation

We introduce the constraint into the objective with Lagrangian relaxation:

$$\begin{aligned} g(\lambda) &= \max_{p(u) \in \{1, \dots, P\}} \sum_u y_{u,p(u)}^{\text{rev}} - \lambda \left(C_{ROAS} \sum_u y_{u,p(u)}^{\text{cost}} - \sum_u y_{u,p(u)}^{\text{rev}} \right) \\ &= \sum_u \max_{p(u) \in \{1, \dots, P\}} \left[(1 + \lambda) y_{u,p}^{\text{rev}} - \lambda C_{ROAS} y_{u,p}^{\text{cost}} \right]. \end{aligned}$$

For any fixed $\lambda = \lambda'$, we can get the optimal $p(u)$ simply by finding the maximum among P discrete values for each u , i.e.,

$$p_{\lambda'}(u) = \operatorname{argmax}_p \{ (1 + \lambda') y_{u,p}^{\text{rev}} - \lambda' C_{ROAS} y_{u,p}^{\text{cost}} \}, \quad \forall u.$$

Then, we can compute the subgradient with respect to λ ,

$$\partial g(\lambda)|_{\lambda=\lambda'} = \sum_u y_{u,p_{\lambda'}(u)}^{\text{rev}} - C_{ROAS} \sum_u y_{u,p_{\lambda'}(u)}^{\text{cost}}$$

and use the sign of the subgradient to guide where to find λ that minimize $g(\lambda)$.

Algorithm 1 Bidding Optimization with Lagrangian Relaxation

Input: $y_{u,p}^{\text{rev}}, y_{u,p}^{\text{cost}}, C_{ROAS}, \lambda_{\text{init}} > 0$, and $0 < \epsilon \ll \lambda_{\text{init}}$

Output: $p^*(u)$ for each u

```

set  $\lambda_{\min} \leftarrow 0$  and  $\lambda_{\max} \leftarrow \lambda_{\text{init}}$ 
if  $\partial g(\lambda)|_{\lambda=0} \geq 0$  then
    return  $p_0(u)$ 
else
    while  $\partial g(\lambda)|_{\lambda=\lambda_{\max}} < 0$  do
        set  $\lambda_{\min} \leftarrow \lambda_{\max}$  and  $\lambda_{\max} \leftarrow 2\lambda_{\max}$ 
    end while
end if
while  $\lambda_{\max} - \lambda_{\min} > \epsilon$  do
    set  $\lambda' \leftarrow \frac{\lambda_{\min} + \lambda_{\max}}{2}$ 
    if  $\partial g(\lambda)|_{\lambda=\lambda'} > 0$  then
        set  $\lambda_{\max} \leftarrow \lambda'$ 
    else if  $\partial g(\lambda)|_{\lambda=\lambda'} < 0$  then
        set  $\lambda_{\min} \leftarrow \lambda'$ 
    else
        break
    end if
end while
return  $p_{\lambda'}(u)$ 

```

A.2 Content Optimization with Online Contextual Bandit

For efficient Bayesian online learning, we will assume the priors for all weights are independent and Gaussian-distributed. The joint distribution can be written in the following form:

$$p(y, t, s, w_1, w_2, \dots, w_k) = h(y, t) a(t, s) g(w_1, w_2, \dots, w_k, s) \prod_i r_i(w_i),$$

where

$$\begin{aligned} r_i(w_i) &= \mathcal{N}(w_i; \mu_i, \sigma_i^2), \\ g(w_1, \dots, w_K, s) &= \delta(x^T w = s), \\ a(s, t) &= \mathcal{N}(t; s, \beta^2), \\ h(y, t) &= \delta(\rho(t) = y), \end{aligned}$$

and $\delta(\cdot)$ is the indicator function.

In the factor graph, the variable nodes $\{w_i, s, t, y\}$ are connected by the factor nodes $\{r_i, g, a, h\}$. We can calculate the posterior distribution $p(w_i|x, y)$ by message passing in factor graph:

- (1) Consider t as root node.
- (2) Perform message passing from leaf to root and obtain $p(t) = \mathcal{N}(t; \eta, \tau^2) h(y, t)$, where $\eta = \sum_i x_i \mu_i$, and $\tau^2 = \sum_i x_i^2 \sigma_i^2 + \beta^2$.
- (3) If $p(t)$ is not Gaussian, we use moment matching to obtain $\hat{p}(t) = \mathcal{N}(t; \mu_t, \sigma_t^2)$, where $\mu_t = \eta + \Delta$ and $\sigma_t^2 = \tau^2 \gamma$.
- (4) Perform message passing from root to leaf and obtain posterior of model weights,

$$p(w_i) = \mathcal{N}(w_i; \mu_{w_i}, \sigma_{w_i}^2),$$

where,

$$\begin{cases} \sigma_{w_i}^2 &= \sigma_i^2 [1 - x_i^2 \frac{\sigma_i^2}{\tau^2} (1 - \gamma)] \\ \mu_{w_i} &= \mu_i + x_i \frac{\sigma_i^2}{\tau^2} \Delta \end{cases}$$

With the above, we can have following online contextual bandit algorithm that works for high-dimensional settings. The algorithm can be easily parallelized or set up for on-device computing.

Algorithm 2 Content Optimization with Online Contextual Bandit

Input: $\mu_i, \sigma_i^2, \beta^2$

```

for time step  $j = 1, 2, \dots, J$  do
    Sample  $\hat{w}_i$  from  $N(\mu_{w_i}, \sigma_{w_i}^2)$  for all  $i$ ;
    for creative  $a = 1, 2, \dots, A$  do
        Sample  $\hat{\epsilon}_a$  from  $N(0, \beta^2)$ 
        Calculate latent reward  $\tilde{t}_{j,a} = \hat{\epsilon}_a + \sum_i x_{i,j,a} \hat{w}_i$ 
    end for
    Choose creative  $a_j = \operatorname{argmax}_a \tilde{t}_{j,a}$  and show it to user;
    observe user engagement and receive reward  $y_j$ 

    Update  $\mu_{w_i} \leftarrow \mu_{w_i} + x_i \frac{\sigma_{w_i}^2}{\tau^2} \Delta$ ;
    Update  $\sigma_{w_i}^2 \leftarrow \sigma_{w_i}^2 [1 - x_i^2 \frac{\sigma_{w_i}^2}{\tau^2} (1 - \gamma)]$ 
end for

```

Here, the value of γ and Δ depends on the link function. For linear regression, $\gamma = 0$ and $\Delta = y - x^T w$. For probit regression, $\gamma = 1 - \omega(\frac{y\eta}{\tau})$ and $\Delta = y\tau\vartheta(\frac{y\eta}{\tau})$, where, $\vartheta(t) = \frac{N(t;0,1)}{\Phi(t;0,1)}$ and $\omega(t) = \vartheta(t)(\vartheta(t) + t)$.

A.3 Sampling from Primal Solution

The IP to LP relaxation can be interpreted as transforming the deterministic integer decision to the expected value of a random integer decision, i.e., the probability of sending some marketing email. We can add randomization after the LP solution to "recover"

the random integer decision $a'_{u,j}$. Formally, we want to implement the sampling so that:

$$E(a'_{u,j}) \propto a_{u,j} \\ \sum_j a'_{u,j} \leq C_{\text{fcap}}.$$

We can follow the following algorithm to "recover" the random integer decision $a'_{u,j}$:

- (1) Calculate sampling probabilities $p_{u,j} = \frac{a_{u,j}}{\sum_j a_{u,j}}$
- (2) Sample a set \mathcal{J}'_u of size $\lfloor \sum_j a_{u,j} + 0.5 \rfloor$ without replacement, from set of all index \mathcal{J}_u with non-uniform distribution of $p_{u,j}$
- (3) We obtain $a'_{u,j} = I(j \in \mathcal{J}'_u)$, where $I(\cdot)$ is indicator function.

B ADAPTIVE HEURISTICS

There are two main classes of AH: 1) Feedback Loop Controller and 2) Hierarchical Optimization. At a high-level, these two classes of AH help optimize and stabilize the system on the "time" and "space" dimensions respectively.

Besides the two main classes of AHs, there are other simple AHs, for example: cool-off rules so that a member cannot receive more than some number marketing messages over a fixed time window, and ϵ -greedy exploration so that selection bias is measured and taken into account in modeling. Most of these heuristics are common to RecSys, and we refer readers to standard practice in [8, 10].




Figure 4: Feedback Loop Controller.

B.1 Feedback Loop Controller

Controllers can observe the feedback from the marketing units, and is able to adjust input to the main optimization module, so as to stabilize it when there are deviations on the constraints. This is particularly useful when we have strict constraints, e.g., marketing budgets/cost. Technically, this can be achieved either by a Proportional–integral–derivative (PID) controller and/or a lightweight reinforcement learning (RL) model. Figure 4 shows the overall flow.

Without loss of generality, assume there is a marketing cost metric C_1 we want to monitor and control from the constraint formulation (2) in Section 2.2.2. The controller takes observed cost C'_1 , and other historical data x^H , and outputs the cost target C_1^* which should be used as the budget in the linear program, i.e.,

$$C_1^* = g(C_1, C'_1, \phi(x^H)).$$

For the case of a PID controller, $\phi(x^H)$ represents the three quantities, proportion, integral, and derivative, which are derived from the time series of error term $e_t(x^H) = C'_{1,t} - C_{1,t}$. Once C_1^* is obtained, we can use it for main optimization module:

$$\begin{aligned} \max_{a_u} \quad & \sum_u f_0(x_u, x_{a_u}, x_c), \\ \text{s.t.} \quad & \sum_u f_1(x_u, x_{a_u}, x_c) \leq C_1^*, \\ & \sum_u f_k(x_u, x_{a_u}, x_c) \leq C_k, \quad \forall 2 \leq k \leq K, \end{aligned}$$

Beside PID controllers, we can also use RL to learn more complex $g(\cdot)$ and $\phi(\cdot)$ functions.

B.2 Hierarchical Optimization

The formulation of Hierarchical Optimization is similar to the LP formulations in Section 2.2.2, except that 1) we use aggregated data as input, 2) the prediction model can be light-weight, and 3) the optimization problem is small-scale.

The main use case for hierarchical optimization is macro-level marketing budget allocation. With slight abuse of notation from Section 2.2.2, we let u index unit of budget allocation (e.g., combination of marketing channel and lines of business), and let a_u denote the level of budget to be assigned to unit u . The resulting optimization formulation is the same as (1) and (2).

The model used in the objective is typically the revenue on budget changes $y^0 = f_0(x_u, x_{a_u}, x_c)$, where either traditional Media Mixed Models (MMM) or modern causal ML methods can be used. Among the constraints, the cost model $y^1 = f_1(x_u, x_{a_u}, x_c)$ can either use learned or rule-based $f_1(\cdot)$ (e.g., assuming cost equals fixed proportion of budget) depending on the maturity of the optimization engine. The constrained optimization problem is small-scale, so we can typically use off-the-shelf solvers to solve the original integer program without LP relaxation.

C ROBUSTNESS OF T-TEST

C.1 Zero-Inflated Heavy-Tail Metric

We can see that in Figure 6, there is a significant probability mass at zero and right side of the distribution is heavy-tailed for Surrogate Index. Note that the counts on histogram plot is \log_{10} scale.




Figure 6: Distribution of Surrogate Index.

C.2 Simulation Result for t -Test Robustness

We simulate the control and treatment based on zero-inflated heavy-tail data. The control is simulated by $X = Z_X X'$, where Z_X is random

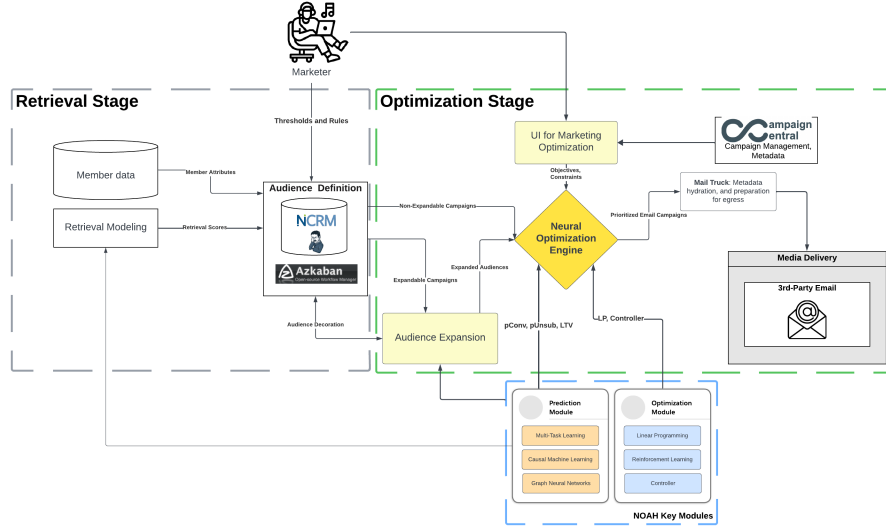


Figure 5: Two-stage Email Marketing System with NOAH

binary variable with $\mathbb{E}[Z_x] = p_0$, and X' is sampled from a heavy-tail distribution. The treatment is simulated by $Y = Z_y Y'$, where $\mathbb{E}[Z_y] = p_0 + p_\Delta$ and $\mathbb{E}[\log(Y')] = \mathbb{E}[\log(X')] + \Delta$.

Table 3: Type I errors of t test against zero-inflated heavy-tail distribution.

α^1	$X' \sim P_{\text{obs}}$		$X' \sim P_{\text{ln}}$	
	$p_0 = 0.02$	$p_0 = 0.8$	$p_0 = 0.02$	$p_0 = 0.8$
0.05	0.03038	0.04837	0.02464	0.03917
0.01	0.00227	0.00858	0.00168	0.00524
0.005	0.00072	0.00376	0.00048	0.00211
0.001	0.00002	0.0005	0.00002	0.00031

¹ Size of the test.

Table 4: Power of t test against zero-inflated heavy-tail distribution.

p_Δ	Δ	$X' \sim P_{\text{obs}}$	$X' \sim P_{\text{ln}}$
0	1.0	0.451	0.153
0	2.0	0.909	0.496
0.005	1.0	0.673	0.295
0.005	2.0	0.976	0.637

The null model is simulated by setting $\Delta = 0$ and $p_\Delta = 0$. Besides varying the size of the test (α), we also (i) vary p_0 at 0.02 and at 0.8, (ii) vary distribution of X' at P_{obs} and P_{ln} . Here P_{obs} is the empirical distribution of the positive Surrogate Index in the A/B test, and P_{ln} is the lognormal distribution with $\mu(\log(X')) = 2$ and $\sigma(\log(X')) = 2$. We perform t-test on the simulated samples and

calculated the estimated Type I errors. We can see Type I errors are well controlled as shown in Table 3.

The alternative model is simulated by varying the values of Δ and p_Δ . For simplicity, we fix $\alpha = 0.05$, $p_0 = 0.2$ and vary the distribution of X' at P_{obs} and P_{ln} . We can see in Table 4 that power increases as p_Δ increases or as Δ increases.

D ADDITIONAL A/B TEST RESULTS

D.1 Distribution of Long-Term Metrics on 2B Conversions




Figure 7: Distribution of Surrogate Index for 2B Conversions.




Figure 8: Distribution of 3-Month Metric for 2B conversions.